

Targeted Failure-State Initialization for On-Policy Distillation: An Audited Case Study in a Persistent Tool-Use Agent

Anonymous ACL submission

Abstract

On-policy distillation supervises only states visited by the student. This is restrictive for persistent agents: early failures can prevent them from reaching states at which a teacher is more capable. We study a 2B-parameter language model acting for six hours through 17 typed tools in an open-source MMORPG. In one reported run per arm, natural-visitation distillation changes tool use and interaction tempo but leaves the quest score at 12/30. A second round augments student rollouts with verified intermediate world states. Evaluated from a fresh, unseeded world, the resulting policy scores 15/30; all three clustered prompt variants pass a wall that none passes under the reported base or first-round policy. This is a weights-only observation, not a replicated causal estimate: the training rounds differ in more than state source. A targeted teacher-forcing probe also finds that malformed history can locally reverse the teacher’s syntax preference. In rewritten runtime logs, correction markers do not repeat within a session. These observations motivate a confirmatory comparison against natural OPD and a success-prefix temporal curriculum, with model weights and interface recovery isolated. We present an audited case study and falsifiable protocol, not established cross-task generality.

1 Introduction

Post-training a small language model to act over long horizons is not just a token-prediction problem. The model chooses which environment states become training data. If it never reaches a state where its teacher knows what to do, dense teacher feedback is unavailable exactly where it is needed. This creates a state-visitation bottleneck: improving token-level agreement on common states can alter

style without transferring the competence that distinguishes the teacher.

We examine that bottleneck in a persistent tool-use setting. A Qwen-family 2B student plays a tile-based MMORPG through a typed action interface. The world persists across conversation rollovers, and benchmark progress requires long chains of gathering, navigation, dialogue, and quest actions. A scaffolded 4B teacher reaches states that the 2B student does not. We apply reverse-KL on-policy distillation to the student’s own emitted tokens and compare natural student visitation with a mixture that begins some rollouts from verified intermediate database states.

The present evidence is hypothesis-generating and under-replicated. It contains one complete six-hour run per principal arm; the three agents within a run are prompt variants sharing weights, harness, and launch conditions, not independent experimental units. We therefore separate observations from claims throughout. The observed sequence is: base 12/30, natural-visitation round 12/30, and state-augmented round 15/30. In the last arm, every prompt variant passes the same wall from an unseeded evaluation world. This supports the hypothesis that training visitation matters, but it does not isolate state seeding from all round-to-round changes or estimate population-level uncertainty.

This paper makes three bounded contributions:

1. We instantiate targeted failure-state reset as a control point for on-policy distillation and document an exploratory case in which a seeded training round is followed by unseeded passage of a previously unseen wall.
2. We distinguish behavioral imitation from

competence: the first round closely changes measured tool mix and tempo toward the teacher while the benchmark score remains fixed.

3. We expose an interface failure that token-average analyses obscure: in a targeted teacher-forcing probe, malformed history reverses the teacher’s local preference. We separate this grading-side observation from a runtime recovery result and specify the experiments required to test generality.

The broader lesson is methodological. Model weights, visited states, and runtime affordances are different interventions. A credible agent-training result must preserve and cross them rather than roll them into one headline.

2 Setting

2.1 Persistent tool-use benchmark

The environment is Kaetram, an open-source tile-based MMORPG backed by a game server and MongoDB. Each policy controls three agents with different priority prompts (grinder, completionist, and explorer) through 17 typed tools. Tools cover observation, navigation, combat, gathering, dialogue, quest inspection, inventory use, and crafting. The prompt includes procedural quest guidance, so the benchmark measures execution of a supplied plan rather than autonomous world-model discovery.

The primary development metric, CORE-3, counts newly reached stages in three quest chains: Forestry (3 stages), Herbalist’s Desperation (3), and Rick’s Roll (4). Each agent can gain at most ten stages; a three-agent run therefore scores in $[0, 30]$. Runs begin from a freshly initialized world and continue for six hours. Conversation sessions roll over under a fixed context budget while game state persists. The intended confirmatory experimental unit is a full run from a fresh world. Prompt variants are clustered observations within that run.

2.2 Agent and interface contract

The policy acts in an observe–reason–act loop through an OpenAI-compatible model endpoint and a typed tool server. Navigation and selected interactions contain deterministic affordances such as path finding and auto-

matic walking; the model still decides goals, tools, and arguments. The current development record contains an important interface asymmetry: older base-model runs used a native model-visible tool schema while some fine-tuned runs did not receive an identical serialization. We therefore do not use those comparisons as primary evidence. The confirmatory implementation requires normalized full-schema hash parity across training and every evaluation arm.

3 On-Policy Distillation Under Visitation Constraints

Let π_S and π_T denote student and teacher policies, and let d^{π_S} be the serving-state distribution induced by student rollouts. A reverse-KL on-policy objective can be written

$$\mathcal{L}_{\text{OPD}} = \mathbb{E}_{s \sim d^{\pi_S}} \mathbb{E}_{a \sim \pi_S(\cdot | s)} [\log \pi_S(a | s) - \log \pi_T(a | s)]. \quad (1)$$

The teacher scores the student’s emitted tokens. Training uses a clipped importance-weighted update around the rollout policy, with advantages frozen at data-build time. Full implementation details appear in Appendix A.

Equation 1 makes the bottleneck explicit. If a teacher advantage is concentrated on a set of states H for which $\Pr_{s \sim d^{\pi_S}}[s \in H] \approx 0$, token-dense grading supplies almost no information about that advantage. We change the state source, not the loss: some student rollouts begin from verified intermediate environment states, and their records are mixed with natural student rollouts. The student still acts and the same teacher still grades every emitted token.

The seeded states are typed game states rather than textual demonstrations. Player rows specify quest stage, skill level, inventory, and a verified walkable position. End-to-end checks confirm that the game loads each state and exposes the intended condition through the same tools used at evaluation. Evaluation is always unseeded. This resembles reset-based curricula in reinforcement learning (Salimans and Chen, 2018; Resnick et al., 2018; Ecoffet et al., 2021). TCOOD uses teacher-success prefixes to initialize multi-turn OPD curricula (Wang et al., 2026). Our narrower target is a verified failure state selected for high

student-visitation deficit and teacher competence, evaluated against both natural OPD and a TCOD-style success-prefix baseline.

4 Evidence Protocol

We grade evidence before interpreting it. A *run* is the independent unit; agents within a run share the policy, launch, infrastructure, and analysis pipeline. Exact scores from single runs are descriptive. Counts over agents or sessions diagnose mechanisms but cannot be treated as independent replications of a policy effect.

The principal comparison uses the same 2B base checkpoint, six-hour unseeded evaluation protocol, benchmark, and nominal harness. Round 1 trains on natural base-policy rollouts. Round 2 initializes from round 1 and mixes natural round-1 rollouts with 2,326 records collected by the round-1 policy from a verified Herbalist milestone state. The round-2 corpus contains 7,849 states, split into 7,024 training and 825 held-out records. Because round 2 also changes the initialization and data batch, round 1 versus round 2 is not a clean seeded-versus-unseeded training ablation.

The *wall* column measures passage of Herbalist stage 1 by the three prompt variants. It is a useful within-run consistency check, not a binomial sample of independent policies. We preserve raw pre-rewrite emissions for confirmatory format analyses; rewritten history is not an authoritative defect counter.

5 Exploratory Results

5.1 Natural-visitation OPD transfers behavior, not score

Round 1 leaves CORE-3 unchanged at 12/30 and every agent reaches the same 4/10 stage profile as base. Nevertheless, several behavioral measurements move toward the 4B teacher (Table 2). The average number of turns per session exactly matches the teacher, observation and navigation shares approach the teacher, and query frequency overshoots it. The result is consistent with mode-seeking behavioral imitation, while showing that agreement on frequently visited states is insufficient for this quest wall.

Per-verb reverse-KL signs and behavioral movement are also aligned in the development

logs: verbs whose held-out divergence falls generally move toward the teacher’s usage, while query-quest both diverges and behaviorally overshoots the teacher. This is descriptive triangulation. It does not establish that a specific token update caused downstream quest behavior.

5.2 The seeded-training round is followed by unseeded wall passage

The 4B teacher demonstrates competence beyond the Herbalist stage-1 wall, while the base 2B almost never visits the required state naturally. Seeded collection starts the round-1 policy at the accepted quest with the required gathering skill and near the target resource. All three seeded collectors pass the wall, creating student-authored actions and teacher scores in this region.

After training on the natural-plus-seeded mixture, round 2 is evaluated from a fresh unseeded world with recovery disabled. It scores 15/30, with each prompt variant at 5/10. All three pass Herbalist stage 1, compared with zero of three under both the base and round-1 policies. Logs show earlier quest acceptance, explicit tracking of the skill gate, resource rotation after empty harvests, and less premature return to the quest giver. These are concrete trajectory differences consistent with the intended mechanism.

The valid statement is deliberately narrow: *a seeded-training round was followed by unanimous unseeded wall passage in one run*. The present data do not show that seeding alone caused the improvement. The required causal experiment trains fresh students from the same checkpoint under identical budgets, seeds, and data construction, varying only whether the verified state mixture is included (Section 7).

5.3 Weights and recovery solve different failures

Round 3 combines broader state seeding with a runtime recovery affordance and scores 18/30. That number is not a pure-weights result. Applying the same recovery affordance to round-2 weights produces 17/30 in one run, suggesting that most of the round-2-to-round-3 stage difference may come from the interface. The clean weights-only comparison remains base 12/30 versus round 2 15/30, one run per arm.

Table 1: Observed development runs and the claims they permit. Scores are descriptive because there is one run per arm. Recovery is a runtime interface affordance.

Policy	Training states	Recovery	CORE-3	Wall	Interpretation
Base 2B	–	off	12	0/3	Reference run
OPD round 1	natural	off	12	0/3	Behavior changed; score did not
OPD round 2	natural + seeded	off	15	3/3	Weights-only, unseeded observation
Round-2 weights	–	on	17	–	Recovery ablation; one run
Round-3 weights	broadly seeded	on	18	–	Weights plus interface; not pure weights

Table 2: Behavior changes in the first OPD round. Percentages are shares of tool calls in one run per policy; they are diagnostics, not independent trials.

Measure	Base	Round 1	Teacher
Turns per session	9.1	6.1	6.1
Observe share	26.6%	35.5%	36.1%
Navigate share	14.3%	9.4%	9.1%
Attack share	17.9%	6.6%	9.1%
Query-quest share	5.3%	18.0%	8.7%
Chars./turn	411	1,248	864

The recovery mechanism detects a dropped malformed call, reconstructs the executable call, rewrites the retained history into canonical form, executes the action, and returns an explicit correction note. In the audited round, 405 sessions contain a recovery marker; each contains exactly one such marker and no logged recurrence after the correction. This is a system-counting fact. It does not show that weights learned the format or that recovery generalizes to other defects.

6 A Teacher-Forcing Failure at the Interface

Round 1 introduces a malformed parameter-key pattern that is absent in the base run. The emitted form places a Boolean value inside the parameter key, so the parser silently drops the intended argument. The training records contain the correct syntax, and the teacher favors the correct delimiter in clean base-policy contexts. A targeted two-sided probe, however, gives a different result after the malformed prefix is already present: at the divergence token, the 4B teacher assigns log probability -0.161 (about 85% probability) to continuing the malformed form, versus -0.253 under the student. On a matched correct-form state, the preference favors the correct continuation.

Thus teacher generation and teacher-forced grading are distinct objects. A teacher that

generates canonical syntax from a clean history can locally reward continuation of a student’s defect after that defect enters the context. Under dense on-policy distillation, this creates wrong-signed feedback at precisely the visited malformed prefix. We call this a *teacher-forcing copy-prior hypothesis*. The current evidence is one targeted defect family and a small probe, so the name describes a candidate mechanism rather than a general law.

This observation relates to exposure bias and to failures of token-level teacher agreement on degraded prefixes. It also gives a testable prediction: paired histories differing only in canonical versus malformed prior syntax should change teacher log-odds while holding state and candidate continuations fixed. Section 7 turns that prediction into a multi-state, multi-defect experiment.

7 Confirmatory Study Design

The repository contains launch and validation work for the following design; results are not reported until immutable run bundles exist.

Primary endpoint and unit. The independent unit is a six-hour run from a fresh seeded world. The preregistered primary endpoint is total CORE-3 stage gain across the three clustered agents. Quest-wall passage is secondary. All runs, confidence intervals, effect sizes, and stopping rules must be reported.

Visitation ablation. Train two fresh students from the same checkpoint with every setting fixed: natural student visitation only versus the same natural data budget mixed with verified milestone states. Match teacher, optimization, action-token budget, environment interactions, and training-seed schedule; evaluate both unseeded.

Weights by recovery factorial. Evaluate base, round-2, and round-3 weights with recov-

ery both off and on. Preserve raw pre-rewrite emissions. This estimates main effects and interactions for stage progress and format-defect rate instead of attributing a combined system score to weights.

Generalization. Hold out at least one quest from state seeding and compare full walk-through against no-walkthrough prompts. A second environment or model family is needed for any broad claim. Plain natural-visitation OPD, corrected-interface SFT, and a matched TCOD-style success-prefix curriculum are minimum baselines under matched interaction and teacher-scoring budgets.

Copy-prior test. Construct paired canonical/malformed histories across tools, states, defect types, and at least two teacher sizes or families. Score identical canonical and malformed continuations; report teacher and student log-odds with paired confidence intervals. Then intervene on grading context and test generation, rather than inferring generation from grader-side movement.

8 Related Work

GKD establishes on-policy distillation from student-generated sequences with teacher feedback (Agarwal et al., 2024); our work does not claim novelty for OPD itself. TCOD directly studies temporal curricula and intermediate-state initialization for multi-turn OPD (Wang et al., 2026); broad novelty for seeded starts is therefore preempted. DAgger formalizes iterative data collection under the learner’s state distribution (Ross et al., 2011). Reset-along-demonstration, Backplay, and Go-Explore show that restoring intermediate states can expose useful signal in hard-exploration reinforcement learning (Salimans and Chen, 2018; Resnick et al., 2018; Ecoffet et al., 2021). We use the same control point for on-policy language-model distillation in a persistent, typed environment, but the current study does not establish superiority over those curriculum families.

Game-agent benchmarks already cover long-horizon language and vision agents. BALROG emphasizes knowing–doing gaps across games (Paglieri et al., 2025), while Orak uses MCP across twelve games and releases gameplay fine-tuning data (Park et al., 2026). MCP and

game play are therefore infrastructure here, not novelty claims. Voyager and AgentTrek use curricula, reusable skills, or guided replay to obtain long-horizon behavior (Wang et al., 2024; Xu et al., 2025). Our narrower question is whether state visitation controls which teacher capabilities are available to a small student under dense grading.

Supervised imitation can compound errors under distribution shift (Ross et al., 2011), and knowledge distillation does not guarantee functional agreement between teacher and student (Stanton et al., 2021). The r10 cross-vocabulary SFT regression in our development record is motivating evidence, but its raw artifact bundle is incomplete and it is not a primary result of this submission draft.

9 Conclusion

The observed case supports a concrete hypothesis: for persistent tool-using agents, controlling where the student acts may matter as much as optimizing how it matches a teacher. Natural-visitation OPD changed behavior without changing score; a later state-augmented training round was followed by unseeded passage of a previously unseen wall. A malformed-prefix probe further shows why teacher-forced token feedback can fail at the tool boundary. Replication, isolation, and held-out transfer are the next scientific results, not boxes that prose can check.

Limitations

The central limitation is statistical: one run per principal arm cannot support a population claim, and three prompt variants are not independent seeds. The state-seeded round also differs from the preceding round in initialization and data. The environment is a single game with procedural walkthroughs, the teacher and student are from one model family, and stage counts are coarse. Some development analyses depend on logs whose raw inputs are not yet packaged. The paper must not be submitted as confirmatory until the study design above is run and a clean clone regenerates every main table.

Direct database construction is environment-specific and may not transfer to systems without restorable state. The failure state was

selected retrospectively, so selection bias is possible. Teacher competence at the target state is observed in a small number of trajectories, not estimated precisely. The syntax probe covers one defect family and should not be generalized to all teacher-forced grading.

Ethical Considerations

The work concerns game agents acting in isolated local worlds. Risks include overstating autonomy, silently baking procedural knowledge into prompts, and mistaking recovered actions for model competence. We mitigate these by labeling scaffolding, logging model-visible schemas and raw emissions, separating runtime recovery from weights, and reporting failed interventions. No human-subject or private-user data are used.

Acknowledgments and AI Assistance Disclosure

Large language models materially assisted literature search, code review, experimental-design critique, and manuscript drafting. Human authors are responsible for every claim, citation, artifact, and submission decision; the scope will also be disclosed in the Responsible NLP Checklist.

References

Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos Garea, Matthieu Geist, and Olivier Bachem. 2024. On-policy distillation of language models: Learning from self-generated mistakes. In *International Conference on Learning Representations*.

Adrien Ecoffet, Joost Huizinga, Joel Lehman, Kenneth O. Stanley, and Jeff Clune. 2021. First return, then explore. *Nature*, 590:580–586.

Davide Paglieri, Bartłomiej Cupiał, Samuel Coward, Ulyana Piterbarg, Maciej Wołczyk, Akbir Khan, Eduardo Pignatelli, Lukasz Kuciński, Llerel Pinto, Rob Fergus, Jakob Foerster, Jack Parker-Holder, and Tim Rocktäschel. 2025. Balrog: Benchmarking agentic llm and vlm reasoning on games. In *International Conference on Learning Representations*.

Dongmin Park, Minkyu Kim, Beongjun Choi, Junhyuck Kim, Keon Lee, Jonghyun Lee, Inkyu Park, Byeong-Uk Lee, Jaeyoung Hwang, Jaewoo Ahn, Ameya S. Mahabaleshwarkar, Bilal Kartal, Pritam Biswas, Yoshi Suhara, Kangwook

Lee, and Jaewoong Cho. 2026. Orak: A foundational benchmark for training and evaluating llm agents on diverse video games. In *International Conference on Learning Representations*.

Cinjon Resnick, Robyn Eldridge, David Ha, Denny Britz, Jakob Foerster, Julian Togelius, Kyunghyun Cho, and Joan Bruna. 2018. Backplay: “man muss immer umkehren”. In *NeurIPS Deep Reinforcement Learning Workshop*.

Stéphane Ross, Geoffrey Gordon, and Drew Bernstein. 2011. A reduction of imitation learning and structured prediction to no-regret online learning. In *International Conference on Artificial Intelligence and Statistics*.

Tim Salimans and Richard Chen. 2018. Learning montezuma’s revenge from a single demonstration. *arXiv preprint arXiv:1812.03381*.

Samuel Stanton, Pavel Izmailov, Polina Kirichenko, Alexander A. Alemi, and Andrew Gordon Wilson. 2021. Does knowledge distillation really work? In *Advances in Neural Information Processing Systems*.

Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. 2024. Voyager: An open-ended embodied agent with large language models. *Transactions on Machine Learning Research*.

Jiaqi Wang, Wenhao Zhang, Weijie Shi, Yaliang Li, and James Cheng. 2026. Tcod: Exploring temporal curriculum in on-policy distillation for multi-turn autonomous agents. *arXiv preprint arXiv:2604.24005*.

Yiheng Xu, Dunjie Lu, Zhennan Shen, Junli Wang, Zekun Wang, Yuchen Mao, Caiming Xiong, and Tao Yu. 2025. Agenttrek: Agent trajectory synthesis via guiding replay with web tutorials. In *International Conference on Learning Representations*.

A Training Details

The teacher is a scaffolded Qwen-family 4B model and the student begins from a 2B checkpoint. Both use the same vocabulary. Student rollouts are reconstructed as the byte-exact serving prompt plus the verbatim emitted continuation. This prefix property avoids re-rendering differences caused by chat-template handling of reasoning delimiters. Teacher and behavior log probabilities are computed on the identical raw string. Every tenth session is held out for development gates.

The update uses PPO-style clipped importance weighting with clip $\epsilon = 0.3$, advantage clamp $[-3, 3]$, and a 1.5 weight on the first

third of each session. Training uses one epoch, learning rate 5×10^{-5} , effective batch size 32, and a fresh LoRA initialized over the prior merged checkpoint. Action positions alone receive the language-model head. These choices were development decisions, not independently ablated.

Round 1 contains 5,564 train and 574 held-out records from the base policy’s natural roll-outs. Round 2 combines 5,523 natural round-1 records and 2,326 seeded records before the 7,024/825 train/held-out split. The seed encodes accepted Herbalist stage 1, Foraging level 5, starter inventory, and a verified tile near the target resource. Evaluation never loads this state.

B Claim-to-Evidence Boundary

Round-1 behavior. Descriptive, one run: measured behavior changed while score did not.

State-seeded round. Unreplicated: the seeded-training round was followed by unseeded wall passage in one run.

12 to 15. Least-confounded historical weights-only observation; not an effect estimate, and exact parity remains unavailable.

18/30. Weights-plus-interface system result only.

Copy prior. Context-dependent endorsement reversal for the targeted historical probe.

Recovery. No repeat marker after correction in rewritten logs; raw relapse is unavailable.

Cross-task transfer. Unsupported; no claim.

C Reproducibility Contract

Before confirmatory runs, each immutable run bundle must contain: environment and game revisions; container or dependency locks; model, tokenizer, adapter, and dataset revisions; complete model-visible tool schemas; rendered-prompt hashes; sampling parameters; environment and inference seeds; raw emissions before any rewrite; state-transition logs; analysis outputs; and checksums. Dataset records

must point to source runs and transformations. A clean-clone command must verify hashes and regenerate every table and figure without network-dependent mutable inputs.

The current development artifact does not yet meet this bar. In particular, raw inputs for the historical r10 analysis and some OPD gate summaries are absent, and full dependency installation has not been demonstrated from the repository root. These omissions are release blockers, not appendix footnotes.

D Planned Statistical Analysis

No session- or agent-level pseudo-replication will be used for policy effects. The run is the analysis unit. A preliminary minimum of five runs per arm may be used to estimate run-level variance, followed by a documented power calculation for the confirmatory sample. The primary comparison is two-sided unless a directional alternative and stopping rule are pre-registered before collection. We will report all runs, bootstrap or randomization-based confidence intervals appropriate to the achieved sample, standardized and raw effect sizes, and the full family of secondary endpoints with multiplicity handling. Exact categorical wall-passage counts will be reported descriptively when sample sizes are small.

E Venue and Disclosure Note

This PDF uses the current official ACL review style as an ARR/NAACL 2027 working draft. The NAACL 2027 venue-specific call is not yet published. Recheck the call, page limit, responsible NLP checklist, anonymity, author-review registration, and commitment rules before submission. ARR prohibits simultaneous archival review.