

Auditing a Candidate State-Visitation Bottleneck in On-Policy Distillation: An Agent Case Study and Confirmatory Protocol

Anonymous ACL submission

Abstract

On-policy distillation supervises student-visited states, potentially excluding persistent states where a teacher has an advantage. We audit this failure mode for a 2B agent using 17 tools in a persistent MMORPG. One reported natural-visitation run changes tool use and tempo but not a 12/30 quest score. A later round augments training with rollouts from a hand-selected intermediate player state and is followed by a reported canonical-start score of 15/30; all three clustered prompt variants pass a prior wall. This sequence is descriptive: rounds differ beyond state source, have one run per arm, and lack complete provenance. We specify a matched study that selects training-only states using frozen visitation, teacher advantage, validity, reachability, and recoverability; crosses player state with model-visible history; and compares targeted states with natural visitation, generic resets, successful prefixes, and turn-level guidance. The contribution is an audited case study and protocol, not an established method effect.

1 Introduction

Post-training a small language model to act over long horizons is not just a token-prediction problem. The model chooses which environment states become training data. If it never reaches a state where its teacher knows what to do, dense teacher feedback is unavailable exactly where it is needed. This creates a state-visitation bottleneck: improving token-level agreement on common states can alter style without transferring the competence that distinguishes the teacher.

We examine that bottleneck in a persistent tool-use setting. A Qwen/Qwen3.5-2B student (Qwen Team, 2026) plays a tile-based MMORPG through a typed action interface.

The world persists across conversation rollovers, and benchmark progress requires long chains of gathering, navigation, dialogue, and quest actions. Historical records report a scaffolded Qwen/Qwen3.5-4B teacher run reaching states that the 2B student run does not. We apply reverse-KL on-policy distillation to the student’s own emitted tokens. The historical comparison contrasts natural student visitation with a mixture that begins some rollouts from loadable, walkability-checked intermediate database states; the confirmatory protocol additionally freezes a measurable state-selection rule before training.

The present evidence is hypothesis-generating and under-replicated. It contains one complete six-hour run per principal arm; the three agents within a run are prompt variants sharing weights, harness, and launch conditions, not independent experimental units. We therefore separate observations from claims throughout. The observed sequence is: base 12/30, natural-visitation round 12/30, and state-augmented round 15/30. In the last arm, every prompt variant passes the same wall from a canonical-start evaluation world with no intermediate-state initialization. This motivates testing whether training visitation matters, but it does not isolate state seeding from all round-to-round changes or estimate population-level uncertainty.

This working draft makes two bounded contributions:

1. We audit a state-visitation failure case and separate observations that survive artifact review from causal claims that do not. We provide a confirmatory protocol with controls capable of falsifying the proposed player-state intervention.
2. We distinguish behavioral imitation from

competence in the surviving historical record: the reported post-distillation run moves measured tool mix and tempo toward the teacher while the benchmark score remains fixed.

Separately, one targeted teacher-forcing probe motivates a copy-prior hypothesis at the tool boundary. It does not establish a within-context preference reversal and is not a headline contribution.

The broader lesson is methodological. Model weights, visited states, and runtime affordances are different interventions. A credible agent-training result must preserve and cross them rather than roll them into one headline.

2 Setting

2.1 Persistent tool-use benchmark

The environment is a pinned fork of Kaetram, an open-source tile-based MMORPG (Kaetram contributors, 2026), backed by a game server and MongoDB. Each policy controls three agents with different priority prompts (grinder, completionist, and explorer) through 17 typed tools. Tools cover observation, navigation, combat, gathering, dialogue, quest inspection, inventory use, and crafting. The prompt includes procedural quest guidance, so the benchmark measures execution of a supplied plan rather than autonomous world-model discovery.

The primary development metric, CORE-3, counts newly reached stages in three quest chains: Forestry (3 stages), Herbalist’s Desperation (3), and Rick’s Roll (4). Each agent can gain at most ten stages; a three-agent run therefore scores in $[0, 30]$. Historical notes describe runs as beginning from a freshly initialized player state in a restarted server, but no fail-closed reset receipt or complete shared-world hash survives. Runs continue for six hours while conversation sessions roll over under a fixed context budget and game state persists. The intended confirmatory experimental unit is a full run with an attested reset. Prompt variants are clustered observations within that run.

2.2 Agent and interface contract

The policy acts in an observe–reason–act loop through an OpenAI-compatible model endpoint and a typed tool server. Navigation

and selected interactions contain deterministic affordances such as path finding and automatic walking; the model still decides goals, tools, and arguments. The current development record contains an important interface asymmetry: older base-model runs used a native model-visible tool schema while some fine-tuned runs did not receive an identical serialization. We therefore do not use those comparisons as primary evidence. The confirmatory implementation requires normalized full-schema hash parity across training and every evaluation arm.

3 On-Policy Distillation Under Visitation Constraints

Let $x \in \mathcal{X}$ denote persistent player state (including position, inventory, skills, and quest progress) and $h \in \mathcal{H}$ the model-visible interaction history. The joint occupancy state is their joint value $z = (x, h)$; two rollouts with the same database snapshot but different visible histories need not induce the same action distribution. Let μ initialize player states, $q(h | x)$ initialize visible histories, and $d_{\mu, q}^{\pi}(x, h)$ denote the joint occupancy induced thereafter by policy $\pi(a | h)$ and the environment dynamics, whose transitions depend on the latent player and world state. The model does not observe x directly; it acts only from h , which may contain tool observations derived from x . The canonical serving initializers are μ_0 and q_0 . A reverse-KL on-policy objective can be written

$$\mathcal{L}_{\text{OPD}} = \mathbb{E}_{(x, h) \sim d_{\mu, q}^{\pi_S}} \mathbb{E}_{a \sim \pi_S(\cdot | h)} [\log \pi_S(a | h) - \log \pi_T(a | h)]. \quad (1)$$

The teacher scores the student’s emitted tokens. Training uses a clipped importance-weighted update around the rollout policy, with advantages frozen at data-build time. Full implementation details appear in Appendix A.

Equation 1 makes the bottleneck explicit. If a teacher advantage is concentrated on a set $H \subseteq \mathcal{X} \times \mathcal{H}$ for which $\Pr_{(x, h) \sim d_{\mu_0, q_0}^{\pi_S}} [(x, h) \in H] \approx 0$, token-dense grading supplies almost no information about that advantage. We change the state source, not the loss: some student rollouts begin from intermediate player states, and their records are mixed with natural student rollouts. The student still acts and the same teacher still grades every emitted

token. Because changing x can also change the compatible h , the confirmatory design crosses player-state initialization with matched history construction rather than treating a database snapshot as the complete policy state.

The seeded states are typed persistent player states rather than textual demonstrations. Player rows specify quest stage, skill level, inventory, and a walkability-checked position. End-to-end loadability is necessary but not sufficient for legal reachability. Each confirmatory candidate must also carry either a hash-verified witness trajectory from the canonical start that is replayed by a pinned checker, or an invariant certificate executed by a pinned checker that establishes a valid path to the snapshot. Hash continuity or certificate-shaped metadata alone is not sufficient. Evaluation always starts from the canonical initial player state; prospective runs nevertheless record distinct gameplay-RNG seeds. This resembles reset-based curricula in reinforcement learning (Salimans and Chen, 2018; Resnick et al., 2018; Ecoffet et al., 2021). TCOd uses teacher-success prefixes to initialize multi-turn OPD curricula (Wang et al., 2026); Guided-OPD mixes teacher and student turns and anneals guidance to zero (Li et al., 2026a); ReOPD treats prefix selection as a teacher-reliability problem (Liao et al., 2026). These works preempt generic novelty for intermediate starts or state-distribution design. Our narrower object is a direct, prefix-independent player-state initializer. A candidate state need not lie on a successful teacher trajectory and is selected by a frozen rule rather than by progress alone.

Let $v(x)$ be estimated natural student visitation, and let $p_T(x, q)$ and $p_S(x, q)$ be repeated teacher and student success rates under a frozen history constructor q . The confirmatory selector admits a player state only if it passes player-state and compatibility invariants, is unfinished and task-relevant, is recoverable, satisfies $v(x) \leq \tau_v$, has certified reachability, and satisfies $p_T(x, q) - p_S(x, q) \geq \tau_\Delta$ together with a minimum teacher success threshold. Counts and denominators, source-run identifiers, complete snapshots, and hash-verified validation artifacts are frozen before outcomes. This rule distinguishes the proposed method from the hand-selected milestone used in the exploratory round.

Table 1: Frozen targeted-state admission rule. L and U are 95% Wilson bounds recomputed from hash-bound Boolean trials. Every row must pass.

Criterion	$n_{\min}$	Rule
Visit rate	500	$U(v) \leq .05$
Teacher success	100	$L(p_T) \geq .60$
Teacher-student gap	100 each	$L(p_T) - U(p_S) \geq .30$
Recovery	100	$L(p_R) \geq .80$
Validity/reachability	–	pinned check

4 Evidence Protocol

We grade evidence before interpreting it. For a fixed checkpoint, a fresh-world evaluation run is the independent unit; agents within it share the policy, launch, infrastructure, and analysis pipeline. For a training-method claim, the independent unit is instead a freshly initialized training seed/checkpoint, with evaluation seeds nested within it. Exact scores from single historical runs are descriptive. Counts over agents or sessions diagnose mechanisms but cannot be treated as independent replications of either effect.

The principal historical sequence uses a 2B model lineage, six-hour canonical-start evaluation protocol, benchmark, and nominal harness, but it is not a same-checkpoint comparison. Round 1 trains on natural base-policy rollouts. Round 2 initializes from round 1 and mixes natural round-1 rollouts with 2,326 records collected by the round-1 policy from a loadable, walkability-checked Herbalist milestone state. The round-2 corpus contains 7,849 states, split into 7,024 training and 825 held-out records. Because round 2 also changes the initialization and data batch, round 1 versus round 2 is not a clean state-augmented-versus-natural-only training ablation.

Separately, before observing outcomes we registered a local feasibility pilot using no paid endpoint and crossing Base, round-2, and round-3 public weight snapshots with three paired inference/environment seeds. Each cell uses one completionist prompt identity, recovery off, a canonical tool schema, the same locked tokenizer and patched render contract, and one five-minute canonical-start episode. This pilot tests execution and action throughput only; it is not the six-hour weights-by-recovery factorial and cannot establish policy superiority.




Figure 1: Causal geometry of the proposed matched study. Training arms run in parallel from common initialization; they do not form a sequential curriculum. Model-visible history is crossed explicitly because a database snapshot alone is not the policy state. The fixed-checkpoint weights-by-recovery factorial is a separate diagnostic and cannot estimate this intervention.

Table 2: Observed development runs and the claims they permit. Scores are descriptive because there is one run per arm. Recovery is a runtime interface affordance.

Policy	Training states	Recovery	CORE-3	Wall	Interpretation
Base 2B	–	off	12	0/3	Reference run
OPD round 1	natural	off	12	0/3	Behavior changed; score did not
OPD round 2	natural + seeded	off	15	3/3	Reported recovery-off observation
Round-2 weights	–	on	17	–	Unmatched recovery-on observation
Round-3 weights	broader seeded	on	18	–	Weights plus interface; not pure weights

The *wall* column measures passage of Herbalist stage 1 by the three prompt variants. It is a useful within-run consistency check, not a binomial sample of independent policies. We preserve raw pre-rewrite emissions for confirmatory format analyses; rewritten history is not an authoritative defect counter.

5 Exploratory Results

5.1 Behavior changes without score change in the reported natural-visitation round

Round 1 leaves CORE-3 unchanged at 12/30 and every agent reaches the same 4/10 stage profile as base. Nevertheless, turns per session and several aggregate tool shares move toward the 4B teacher: turns/session are 9.1/6.1/6.1 for base/round 1/teacher, observe shares are 26.6/35.5/36.1%, and navigate shares are 14.3/9.4/9.1%. Query-quest share instead overshoots (5.3/18.0/8.7%). These one-run diagnostics are consistent with behavioral imitation on frequently visited states, but they lack matched-state opportunities and uncertainty. We therefore make no quantitative KL-behavior correspondence or token-level causal claim.

5.2 The seeded-training round is followed by canonical-start wall passage

The 4B teacher demonstrates competence beyond the Herbalist stage-1 wall. In the reported

base and round-1 runs, zero of three prompt variants crosses that wall; the available record does not provide a defensible denominator for a broader natural-visitation rate. Seeded collection starts the round-1 policy at the accepted quest with the required gathering skill and near the target resource. All three seeded collectors pass the wall, creating student-authored actions and teacher scores in this region.

After training on the natural-plus-seeded mixture, round 2 is reported as evaluated from a canonical-start player state in a restarted server, with no intermediate-state initialization and with recovery disabled. Neither a reset receipt nor its historical gameplay-RNG seed was preserved. It scores 15/30, with each prompt variant at 5/10. All three pass Herbalist stage 1, compared with zero of three under both the base and round-1 policies. Logs show earlier quest acceptance, explicit tracking of the skill gate, resource rotation after empty harvests, and less premature return to the quest giver. These are concrete trajectory differences consistent with the intended mechanism.

The valid statement is deliberately narrow: *a seeded-training round was followed by unanimous canonical-start wall passage in one run.* The present data do not show that seeding alone caused the improvement. The required causal experiment trains fresh students from the same checkpoint under identical budgets, seeds, and data construction, varying only whether the legally certified state mixture is

included (Section 7).

5.3 Weights and recovery are historically confounded

Round 3 combines broader state seeding with a runtime recovery affordance and scores 18/30. That number is not a pure-weights result. A separate unmatched run of round-2 weights with recovery reports 17/30. The run directories were recovered externally, but missing configuration parity, checkpoint/render attestation, reset receipts, and gameplay seeds prevent attribution of either score difference to the interface or weights. The reported recovery-off contrast is base 12/30 versus round 2 15/30, one run per arm; exact checkpoint, render, seed, and artifact parity is unavailable, so even this contrast is not a clean causal weights comparison.

The recovery mechanism detects a dropped malformed call, reconstructs the executable call, rewrites the retained history into canonical form, executes the action, and returns an explicit correction note. Historical notes report 405 rewritten sessions with exactly one recovery marker and no later marker. Recovered rewritten bundles exist, but raw pre-rewrite emissions are absent and the 405 count has not been independently regenerated. It therefore does not show that weights learned the format, that the model self-corrected, or that recovery generalizes to other defects.

5.4 Prospective pilot executes all cells, not the task

All nine preregistered local cells completed with exact canonical-start projections and seed-bound environment-RNG receipts. Rehashing the 261 files in their sealed inventories found no mismatch. Across 106 raw generations, 42 emitted a canonical structured call and 64 emitted none, but no cell advanced a quest or a CORE-3 stage. The complete throughput diagnostics appear in Appendix B. The retained bundle predates the v2 database and full-snapshot attestations and is explicitly labeled `legacy_v1_unattested`. This is execution evidence only; five minutes is too short for a task-level or policy-quality estimate.

6 A Teacher-Forcing Failure at the Interface

Round 1 introduces a malformed parameter-key pattern that is absent in the base run. The emitted form places a Boolean value inside the parameter key, so the parser silently drops the intended argument. The training records contain the correct syntax, and historical notes report that the teacher favors the correct delimiter in clean base-policy contexts. After the malformed prefix is already present, the same notes report teacher log probability -0.161 (about 85% probability) for continuing the malformed form, versus -0.253 under the student. This is a positive teacher-over-student distillation advantage for one candidate. The reported correct-form observation uses a different state; without both canonical and malformed candidates scored in each matched context, the record does not establish a preference reversal.

Thus teacher generation and teacher-forced grading are distinct objects, but the historical probe establishes only a candidate failure mode. A teacher that generates canonical syntax from a clean history may still assign a locally useful distillation signal to continuation of a student's defect after that defect enters the context. We call this a *teacher-forcing copy-prior hypothesis*. The current record lacks the paired artifact needed to show that the teacher actively prefers the malformed candidate.

This hypothesis relates to exposure bias (Bengio et al., 2015) and to failures of token-level teacher agreement on degraded prefixes. CCOPD already conditions a teacher on clean, canonical evidence to reduce self-anchored drift in a multi-turn student (Lin et al., 2026); that broad clean-context intervention is not novel here. The narrower tool-syntax hypothesis gives a testable prediction: paired histories differing only in canonical versus malformed prior syntax should change teacher log-odds while holding state and candidate continuations fixed. Section 7 turns that prediction into a multi-state, multi-defect experiment.

7 Confirmatory Study Design

Every matched-training arm below is prospective. The anonymous preparation artifact validates immutable record contracts and common

optimization settings, but it does not contain verified trained-arm bundles or accelerator outcomes. Results will not be reported until every arm, analysis rule, and immutable run bundle passes the same fail-closed protocol.

Only the state-source family and player-state-by-history ablation determine the primary selector claim. The weights-by-recovery study, supervision-structure comparison, and copy-prior probe are separately reported secondary modules; they cannot rescue a failed primary contrast or inflate its arm family.

Primary endpoint and unit. The method estimand is over independently initialized training seeds/checkpoints. Within each training seed, evaluation seeds are nested: each evaluation is a six-hour run from a fresh canonical-start world, with no intermediate player-state initialization and a registered gameplay-RNG seed. The preregistered endpoint aggregates total CORE-3 stage gain across the three clustered agents and then across the frozen evaluation-seed schedule. Quest-wall passage is secondary. The current five-training-seed launcher grid is an unpowered execution template, not a justified sample size; no method claim is permitted until a prospective two-level power analysis freezes both training and nested evaluation replication. All runs, confidence intervals, effect sizes, exclusions, and stopping rules must be reported.

State-source ablation. The six specified primary training contracts are natural OPD, targeted persistent player states, random-valid player states, progress-matched player states, TCOB-B2F teacher-success prefixes, and Guided-OPD. Four prespecified mechanism/baseline arms—visitation-only selection, teacher-advantage-only selection, corrected-interface SFT, and SCoRe-style first-error prefixes—are analyzed separately rather than silently counted as extra primary arms. Match teacher, optimizer, action-token and teacher-scoring budgets, environment interactions, and training-seed schedule; evaluate every arm from the original canonical-start state. A matched Trust-Region Behavior Blending (TRB) arm is additionally required because it changes occupancy through a teacher-near behavior policy while retaining OPD; it is not yet implemented in the launcher. The method claim fails if ran-

dom or progress-matched resets, TCOB-B2F, Guided-OPD, or TRB ties the targeted arm within the preregistered smallest effect of interest.

Supervision structure is a separate axis. Structured Agent Distillation (SAD) uses distinct reasoning- and action-span objectives rather than the historical action-only head (Liu et al., 2026). Before any method-outcome analysis, the study must add and freeze a separately reported natural-visitation comparator that holds trajectories and state source fixed while varying action-only versus SAD-style structured supervision. This arm is not yet registered or implemented; it will not be counted as another state-source arm or conflated with corrected-interface SFT.

Player-state by history ablation. For the same certified player state, compare (i) a direct snapshot with the minimal canonical observation history, (ii) replay of a teacher trajectory with its authentic prefix, and (iii) a direct snapshot with a matched reconstructed history. A fourth Backplay condition anneals starts backward along a witness trajectory. This distinguishes initialization of latent player state x from information or error patterns carried in h , and tests whether a direct snapshot adds value beyond successful prefix replay. The present implementation does not restore a complete shared world (for example NPC, mob, resource, or concurrent-player state), so no full-world-state claim is made.

Weights by recovery factorial. Evaluate base, round-2, and round-3 weights with recovery both off and on. Preserve raw pre-rewrite emissions. These are conditional contrasts among three fixed checkpoint artifacts, not estimates of a training procedure. The registered full study uses 20 paired fresh-world replicate clusters, sums all three personality lanes before comparison, and freezes seven contrasts: round 2 and round 3 versus base with recovery off; recovery on versus off for each checkpoint; and both checkpoint-by-recovery interactions. It controls familywise alpha at .05 across seven paired comparisons. The same simultaneous intervals classify effects against a frozen ± 3 -stage threshold. Under the working normal-model assumption that paired differences have population mean and standard deviation three, 20

clusters give .911 model power for each contrast separately—not for all seven together. This is not evidence about their actual distribution or variability. The separate registered 18-cell local run is exploratory and cannot replace this design; no outcome enters the paper until its complete bundle passes the fail-closed audit. Reporting separates six objects: raw malformed emissions, emissions eligible for correction, harness rewrites or retries, parser-accepted calls, environment-executed calls, and state-changing calls. Harness recovery for a fixed checkpoint is therefore not labeled model self-correction or a learned recovery policy. The running exploratory analysis retains its frozen diagnostics; any additional trace analysis is explicitly post hoc.

Generalization. One quest is now frozen in an evaluation-only v2 registration that supersedes the original task-choice record before any held-out-quest outcome was inspected. The registration hash-binds all aliases, resolved-prompt leakage markers, the Base tokenizer snapshot, and text/token exclusions that every prospective registered training contract must enforce. No verified trained arm bundle yet exists. The no-walkthrough prompt omits game knowledge and is checked after all personality blocks are resolved. We will report reaching the bottleneck, crossing it conditional on arrival, and downstream completion; no held-out outcome is present in this draft. A second environment or model family remains necessary for any broad claim. Corrected-interface SFT and a SCoRe-style first-error-prefix condition are additional baselines under matched interaction and teacher-scoring budgets (Lyu et al., 2025).

Copy-prior test. Construct paired canonical/malformed histories across tools, states, defect types, and at least two teacher sizes or families. Score identical canonical and malformed continuations; report teacher and student log-odds with paired confidence intervals. Then intervene on grading context and test generation, rather than inferring generation from grader-side movement.

8 Related Work

GKD establishes on-policy distillation from student-generated sequences with teacher feed-

back (Agarwal et al., 2024); our work does not claim novelty for OPD itself. Rethinking OPD identifies teacher-student compatibility and genuinely new teacher capability at student-visited states as success conditions (Li et al., 2026c); our historical sequence does not validate a broader visitation principle. SAD changes the unit and weighting of supervision over reasoning and action spans (Liu et al., 2026); our proposed intervention instead changes training-state occupancy while holding the loss fixed. OPCD distills contextual experience on student-generated trajectories, including in text games (Ye et al., 2026); On-Policy Expert Corrections switch from student to expert within a trajectory to address covariate shift (Lauffer et al., 2025). TCOD directly studies temporal curricula and intermediate-state initialization for multi-turn OPD (Wang et al., 2026); broad novelty for seeded starts is therefore preempted. Guided-OPD changes live occupancy by mixing teacher and student turns (Li et al., 2026a); ReOPD designs reliable prefix distributions offline (Liao et al., 2026); TRB anneals a teacher-near rollout policy inside a student-centered trust region (Plyusov et al., 2026); and SCoRe localizes the first student error before short-horizon training (Lyu et al., 2025). TCOD, Guided-OPD, ReOPD, and TRB are concurrent preprints, so we treat them as relevant contemporaneous methods rather than refereed precedent. SOD and KAT study supervision on degraded prefixes (Zhong et al., 2026; Xin et al., 2026). SOD in particular treats post-tool-error high-divergence steps as increasingly unreliable and reweights them; our unvalidated hypothesis is narrower, predicting locally wrong-signed candidate preference. TrOPD likewise filters or changes the divergence direction outside a teacher-verifiable trust region and adds off-policy teacher-prefix guidance (Xing et al., 2026). Unmasking OPD motivates per-token, per-context diagnostics rather than aggregate-loss interpretation (Armandpour et al., 2026). These methods make intermediate-state, student-failure, and state-centric novelty claims unavailable. DAGger formalizes iterative data collection under the learner’s state distribution (Ross et al., 2011). A recent state-distribution analysis likewise argues that the source and locality of training states can matter as much as the loss (Nie,

2026). Reset-along-demonstration, Backplay, and Go-Explore show that restoring intermediate states can expose useful signal in hard-exploration reinforcement learning (Salimans and Chen, 2018; Resnick et al., 2018; Ecoffet et al., 2021). We use the same control point for on-policy language-model distillation in a persistent, typed environment, but the current study does not establish superiority over those curriculum families.

The proposed state selector uses quest stage, inventory, position, and repeated teacher–student success gaps during training. This is privileged information and must be costed and evaluated for held-out transfer. Privileged Information Distillation and its on-policy self-distillation variant already study action-only transfer from privileged agent signals (Penaloza et al., 2026). SERL separately shows that the source and insertion point of environmental feedback matter in multi-turn agents (Li et al., 2026b); neither privilege nor feedback selection is novel here.

Game-agent benchmarks already cover long-horizon language and vision agents. BALROG emphasizes knowing–doing gaps across games (Paglieri et al., 2025), while Orak uses MCP across twelve games and releases gameplay fine-tuning data (Park et al., 2025). Orak emphasizes breadth through modular evaluation across genres, agent/input ablations, data quality, generalization, latency, and a public artifact. Our intended contribution is narrower depth in one persistent game: causal isolation of player state, visible history, weights, and run-time interface. MCP and game play are therefore infrastructure here, not novelty claims. GITM uses LLM-generated plans, text memory, and structured actions in open-world Minecraft (Zhu et al., 2024). Voyager and AgentTrek use curricula, reusable skills, or guided replay to obtain long-horizon behavior (Wang et al., 2024; Xu et al., 2025). Our narrower question is whether a frozen selector over low-occupancy, teacher-competent persistent player states adds value beyond natural OPD, successful-prefix replay, turn-level guidance, and matched generic resets.

Learned tool-error recovery is also distinct from the present runtime factorial. Fission-GRPO converts execution errors into on-policy corrective training instances and evaluates

learned multi-turn recovery (Zhang et al., 2026). Here the recovery-on arm rewrites or retries a fixed checkpoint’s malformed emission in the harness; it cannot establish that the model learned to diagnose or repair its own error.

Supervised imitation can compound errors under distribution shift (Ross et al., 2011), and knowledge distillation does not guarantee functional agreement between teacher and student (Stanton et al., 2021). The r10 cross-vocabulary SFT regression in our development record is motivating evidence, but its raw artifact bundle is incomplete and it is not a primary result of this submission draft.

9 Conclusion

The audited case motivates a concrete hypothesis: low-occupancy player states may hide useful teacher competence from natural OPD. Natural-visitation OPD was followed by behavior differences without a score difference; a later hand-selected state-augmented round was followed by reported canonical-start passage of a wall that the earlier reported policies had not passed. That sequence does not establish the proposed selector. Random and progress-matched resets, TCOD-B2F, Guided-OPD, replication, and held-out transfer decide whether the narrow method survives. A malformed-prefix probe separately motivates, but does not yet establish, a wrong-signed teacher-feedback mechanism at the tool boundary. A prospective local pilot executed all registered cells with internally hash-checked legacy-v1 evidence, but its short horizon produces no quest progress and no policy-quality claim.

Limitations

The central limitation is statistical: one run per principal arm cannot support a population claim, and three prompt variants are not independent seeds. The state-seeded round also differs from the preceding round in initialization and data. The environment is a single game with procedural walkthroughs, the teacher and student are from one model family, and stage counts are coarse. Some development analyses depend on logs whose raw inputs are not yet packaged. The paper must not be submitted as confirmatory until the study design above is

run and a clean clone regenerates every main table.

Direct database construction is environment-specific and may not transfer to systems without restorable state. The historical failure state was selected retrospectively, so selection bias is possible; the prospective selector has no outcome yet. The local pilot has only three five-minute cells per snapshot; it was designed for feasibility, not power, and its outcomes cannot enter the confirmatory estimate. Teacher competence at the target state is observed in a small number of trajectories, not estimated precisely. The syntax probe covers one defect family and should not be generalized to all teacher-forced grading.

Ethical Considerations

The work concerns game agents acting in isolated local worlds. Risks include overstating autonomy, silently baking procedural knowledge into prompts, and mistaking recovered actions for model competence. We mitigate these by labeling scaffolding, logging model-visible schemas and raw emissions, separating runtime recovery from weights, and reporting failed interventions. No human-subject or private-user data are used.

References

Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos Garea, Matthieu Geist, and Olivier Bachem. 2024. On-policy distillation of language models: Learning from self-generated mistakes. In *International Conference on Learning Representations*.

Mohammadreza Armandpour, Fatih Ilhan, David Harrison, Ajay Jaiswal, Duc N. M. Hoang, Farfash Faghri, Yizhe Zhang, Minsik Cho, and Mehrdad Farajtabar. 2026. Unmasking on-policy distillation: Where it helps, where it hurts, and why. *arXiv preprint arXiv:2605.10889*.

Samy Bengio, Oriol Vinyals, Navdeep Jaitly, and Noam Shazeer. 2015. Scheduled sampling for sequence prediction with recurrent neural networks. In *Advances in Neural Information Processing Systems*.

Adrien Ecoffet, Joost Huizinga, Joel Lehman, Kenneth O. Stanley, and Jeff Clune. 2021. First return, then explore. *Nature*, 590:580–586.

Kaetram contributors. 2026. Kaetram-open: An open-source 2d mmorpg. <https://github.com/Kaetram/Kaetram-Open>. Accessed 2026-07-23.

Niklas Lauffer, Xiang Deng, Srivatsa Kundurthy, Brad Kenstler, and Jeff Da. 2025. Imitation learning for multi-turn LM agents via on-policy expert corrections. *arXiv preprint arXiv:2512.14895*.

Gengsheng Li, Mao Zheng, Mingyang Song, Ruiqi Liu, Tianyu Yang, Jie Sun, Qiyong Zhong, Haiyun Guo, Junfeng Fang, Dan Zhang, and Jinqiao Wang. 2026a. On-policy distillation with curriculum turn-level guidance for multi-turn agents. *arXiv preprint arXiv:2606.15912*.

Xiaozhe Li, Tianyi Lyu, Yang Li, Yichuan Ma, Peiji Li, Linyang Li, Qipeng Guo, Dahua Lin, and Kai Chen. 2026b. What and when to distill: Selective hindsight distillation for multi-turn agents. *arXiv preprint arXiv:2605.19447*.

Yaxuan Li, Yuxin Zuo, Bingxiang He, Jinqian Zhang, Chaojun Xiao, Cheng Qian, Tianyu Yu, Huan-ang Gao, Wenkai Yang, Zhiyuan Liu, and Ning Ding. 2026c. Rethinking on-policy distillation of large language models: Phenomenology, mechanism, and recipe. *arXiv preprint arXiv:2604.13016*.

Baohao Liao, Hanze Dong, Christof Monz, Xinxing Xu, Li Dong, and Furu Wei. 2026. Multi-turn on-policy distillation with prefix replay. *arXiv preprint arXiv:2607.04763*.

Zizhuo Lin, Quanling Liu, Jinsheng Quan, Chao Zhang, Yifan Zhu, Xing Shi, Jingtao Xu, Zhihui Li, and Yawei Luo. 2026. Same evidence, different answers: Canonical-context on-policy distillation for multi-turn language models. *arXiv preprint arXiv:2605.30251*.

Jun Liu, Zhenglun Kong, Peiyan Dong, Changdi Yang, Tianqi Li, Hao Tang, Geng Yuan, Wei Niu, Wenbin Zhang, Pu Zhao, Xue Lin, Dong Huang, and Yanzhi Wang. 2026. Structured agent distillation for large language model. In *Proceedings of the 25th International Conference on Autonomous Agents and Multiagent Systems*.

Yuanjie Lyu, Chengyu Wang, Jun Huang, and Tong Xu. 2025. From correction to mastery: Reinforced distillation of large language model agents. *arXiv preprint arXiv:2509.14257*.

Dong Nie. 2026. Post-training is about states, not tokens: A state distribution view of SFT, RL, and on-policy distillation. *arXiv preprint arXiv:2605.22731*.

Davide Paglieri, Bartłomiej Cupiał, Samuel Coward, Ulyana Piterbarg, Maciej Wołczyk, Akbir Khan, Eduardo Pignatelli, Łukasz Kuciński, Lerrel Pinto, Rob Fergus, Jakob Foerster, Jack Parker-Holder, and Tim Rocktäschel. 2025. Balrog: Benchmarking agentic llm and vlm reasoning on games. In *International Conference on Learning Representations*.

845	Dongmin Park, Minkyu Kim, Beongjun Choi, Junhyuck Kim, Keon Lee, Jonghyun Lee, Inkyu Park, Byeong-Uk Lee, Jaeyoung Hwang, Jaewoo Ahn, Ameya S. Mahabaleshwarkar, Bilal Kartal, Pritam Biswas, Yoshi Suhara, Kangwook Lee, and Jaewoong Cho. 2025. Orak: A foundational benchmark for training and evaluating llm agents on diverse video games . <i>arXiv preprint arXiv:2506.03610</i> .	900
846		901
847		902
848		903
849		
850		904
851		905
852		906
853		
854	Emiliano Penalosa, Dheeraj Vattikonda, Nicolas Gontier, Alexandre Lacoste, Laurent Charlin, and Massimo Caccia. 2026. Privileged information distillation for language models. <i>arXiv preprint arXiv:2602.04942</i> .	
855		907
856		908
857		909
858		910
		911
		912
859	Daniil Plyusov, Alexey Gorbatsovski, Alexey Malakhov, Nikita Balagansky, Boris Shaposhnikov, Daria Korotyshova, and Daniil Gavrilov. 2026. Trust-region behavior blending for on-policy distillation. <i>arXiv preprint arXiv:2605.31159</i> .	913
860		914
861		915
862		916
863		
864		
865	Qwen Team. 2026. Qwen3.5-2B and Qwen3.5-4B model cards. https://huggingface.co/collections/Qwen/qwen35 . Accessed 2026-07-23.	917
866		918
867		919
868		920
		921
		922
869	Cinjon Resnick, Roberta Raileanu, Sanyam Kapoor, Alexander Peysakhovich, Kyunghyun Cho, and Joan Bruna. 2018. Backplay: “man muss immer umkehren”. <i>arXiv preprint arXiv:1807.06919</i> .	923
870		924
871		925
872		
873	Stéphane Ross, Geoffrey Gordon, and Drew Bernstein. 2011. A reduction of imitation learning and structured prediction to no-regret online learning. In <i>International Conference on Artificial Intelligence and Statistics</i> .	926
874		927
875		928
876		929
877		930
878	Tim Salimans and Richard Chen. 2018. Learning montezuma’s revenge from a single demonstration. <i>arXiv preprint arXiv:1812.03381</i> .	
879		
880		
881	John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. 2017. Proximal policy optimization algorithms. <i>arXiv preprint arXiv:1707.06347</i> .	931
882		932
883		933
884		934
885		935
886	Samuel Stanton, Pavel Izmailov, Polina Kirichenko, Alexander A. Alemi, and Andrew Gordon Wilson. 2021. Does knowledge distillation really work? In <i>Advances in Neural Information Processing Systems</i> .	936
887		937
888		938
889		
890	Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. 2024. Voyager: An open-ended embodied agent with large language models. <i>Transactions on Machine Learning Research</i> .	939
891		
892		
893		
894		
895	Jiaqi Wang, Wenhao Zhang, Weijie Shi, Yaliang Li, and James Cheng. 2026. Tcod: Exploring temporal curriculum in on-policy distillation for multi-turn autonomous agents. <i>arXiv preprint arXiv:2604.24005</i> .	940
896		941
897		942
898		943
899		944
		945
		946
		947
		948
		949
		950
		951
		952

A Training Details

The settings below are reconstructed from the current implementation and notes written after the historical runs; they are reported settings, not a hash-attested historical training configuration.

The teacher is a scaffolded Qwen/Qwen3.5-4B model and the student begins from Qwen/Qwen3.5-2B (Qwen Team, 2026). Both use the same vocabulary. Student rollouts are reconstructed as the byte-exact serving prompt plus the verbatim emitted continuation. This prefix property avoids re-rendering differences caused by chat-template handling of reasoning

delimiters. Teacher and behavior log probabilities are computed on the identical raw string. Every tenth session is held out for development gates.

The update uses PPO-style clipped importance weighting (Schulman et al., 2017) with clip $\epsilon = 0.3$, advantage clamp $[-3, 3]$, and a 1.5 weight on the first third of each session. Training uses one epoch, learning rate 5×10^{-5} , effective batch size 32, and a fresh LoRA initialized over the prior merged checkpoint. Action positions alone receive the language-model head. These choices were development decisions, not independently ablated.

Round 1 contains 5,564 train and 574 held-out records from the base policy’s natural roll-outs. Round 2 combines 5,523 natural round-1 records and 2,326 seeded records before the 7,024/825 train/held-out split. The seed encodes accepted Herbalist stage 1, Foraging level 5, starter inventory, and a walkability-checked tile near the target resource. Evaluation never loads this state.

B Local Pilot Diagnostics

Table 3: Preregistered local feasibility pilot. Each snapshot has three paired, five-minute cells. Calls/min uses observed wall time, including the final in-flight generation. No-call counts include thinking-only generations.

Weights	Calls/cell	Calls/min	No-call sum
Base	5, 6, 5	0.968	21
Round 2	1, 5, 6	0.736	22
Round 3	5, 4, 5	0.872	21

Every emitted call had a canonical tool name and JSON-object arguments. The retained stderr logs contain no logged **API error:** line, every final session records duration exhaustion, and all sealed inventory hashes reverify. The per-cell call counts overlap substantially; the pilot was not powered for a superiority test and produced no task progress.

C Claim-to-Evidence Boundary

Round-1 behavior. Descriptive, one run: measured behavior changed while score did not.

State-seeded round. Unreplicated: the seeded-training round was followed by

canonical-start wall passage in one run.

Targeted selector. Protocol only: no candidate-state bundle, trained checkpoint, or outcome is currently available.

12 to 15. Reported recovery-off historical contrast; not an effect estimate or clean weights comparison, and exact parity remains unavailable.

18/30. Weights-plus-interface system result only.

Copy prior. Positive teacher-over-student advantage for one malformed continuation in a targeted historical probe; no matched preference reversal.

Recovery. Historical notes report no repeat marker after correction in rewritten sessions. Recovered rewritten bundles exist, but raw pre-rewrite emissions and an independently regenerated marker count are unavailable.

Local pilot. Nine short paired cells executed through the matched local stack with internally hash-checked artifacts; they provide no model-superiority or quest-performance claim and cannot enter the confirmatory estimate.

Historical bundles. The recovered external inventory binds the headline R10 and OPD agent-run directories and reproduces the R10 statistics, but it is not a public reviewer artifact and does not attest the executed reset command, game revision, or OPD training inputs.

Cross-task transfer. Unsupported; no claim.

D Reproducibility Contract

Before confirmatory runs, each immutable run bundle must contain: environment and game revisions; container or dependency locks; model, tokenizer, adapter, and dataset revisions; complete model-visible tool schemas; rendered-prompt hashes; sampling parameters; environment and inference seeds; raw emissions before any rewrite; state-transition logs; analysis outputs; and checksums. Dataset records

must point to source runs and transformations. A clean-clone command must verify hashes and regenerate every table and figure without network-dependent mutable inputs.

The current artifact does not yet meet this bar. An internal recovery hashes 60 selected directories (18,812 files) and regenerates the R10 statistics, but it is not in the reviewer artifact and cannot attest the historical reset command, game revision, OPD training inputs, or render contract. One historical lane also reset a different database from its game servers, so its results are quarantined. The headline R10 and OPD values came from a separate, database-aligned orchestrator and all 36 recovered first states match the canonical level-1 start, but their exact reset command and server revision remain unattested.

A fresh clone now installs and passes the prospective CPU unit suite. Separately, nine matched public-weight pilot cells bind the game revision, tokenizer, render contract, seeds, initial state, and 261-file inventories. This establishes local execution, not training, confirmatory evidence, or clean-clone table regeneration. Neither bundle is reviewer-accessible yet.

Future local launches fail closed on the effective database, mode, configuration, and complete model snapshot. Those protections do not retroactively promote older v1 bundles or the running exploratory factorial. The matched-training contract freezes the arm-seed cells, shared budgets and interfaces, held-out registration, snapshot, tokenizer, and leakage checks. Corrected-interface SFT and first-stage SCoRe-style preparation exist; verified training bundles, SCoRe’s reward stage, live Guided-OPD collection, and accelerator results do not. Release packaging, license review, training, and clean-clone regeneration remain submission blockers.

E Planned Statistical Analysis

No session- or agent-level pseudo-replication will be used. For the causal training comparison, a freshly trained checkpoint is the independent unit and fresh-world evaluation seeds are nested repeated measurements. The primary method contrast must aggregate the frozen nested evaluations within each training seed before uncertainty across training seeds is esti-

mated. The five seeds in the current launcher example are not a power justification.

For the fixed-checkpoint weights-by-recovery study, the independent unit is the paired fresh-world replicate cluster; inference is conditional on the three registered artifacts and cannot estimate training-seed variability. Its seven contrasts, 20 clusters, familywise error rule, and assumption-based power record are frozen prospectively. An executable contract fixes the estimator, intervals, effect scale, zero-variance behavior, and no-rerun rule before outcome access. Completion hash-binds the prelaunch ledger, and analysis requires the same clean registered commit while emitting only portable artifact paths. For future trained-model comparisons, the smallest effect and both replication levels must be set from independent pilot variance or a conservative frozen grid. If variance is unknown, use blinded re-estimation with a fixed maximum size. Primary tests are two-sided unless direction and stopping are preregistered. Report every run and exclusion, raw and standardized primary effects, and design-matched uncertainty. Any inferential secondary family needs its own prespecified multiplicity rule; otherwise it remains descriptive. Small wall-passage counts remain descriptive.